

Exercise 3: Shared Memory and Synchronization

Operating Systems UE 2019W

Michael Platzer

Technische Universität Wien
Computer Engineering
Cyber-physical Systems

2019-12-12

Overview

Inter-process communication

Considered so far...


- ▶ Implicit Synchronization
 - ▶ Blocking read- and write operations
 - ▶ Non-related processes via sockets
 - ▶ Related processes via unnamed pipes

Today...

- ▶ Exchanging data via same memory
 - ▶ Memory Mappings
 - ▶ POSIX Shared Memory (SHM)
- ▶ Explicit synchronization of multiple processes
 - ▶ POSIX Semaphore
 - ▶ Synchronization tasks

Shared Memory

- Common memory area: Multiple processes (related or unrelated) can access the same region in the physical memory (i.e., share data). This memory region is mapped into the address space of these processes.



- Read and modify by normal memory access operations
- Fast inter process communication¹

Concurrent access!

→ Explicit synchronization is necessary

¹no intervention of the OS kernel/“zero-copy”, see

POSIX Shared Memory

Exercise 3: Shared Memory and Synchroniza- tion

M. Platzer

Overview

Shared Memory

POSIX Shared Memory API

Memory Mapping

Example

Semaphores

Motivation

Synchronization Tasks

POSIX Semaphore

Examples

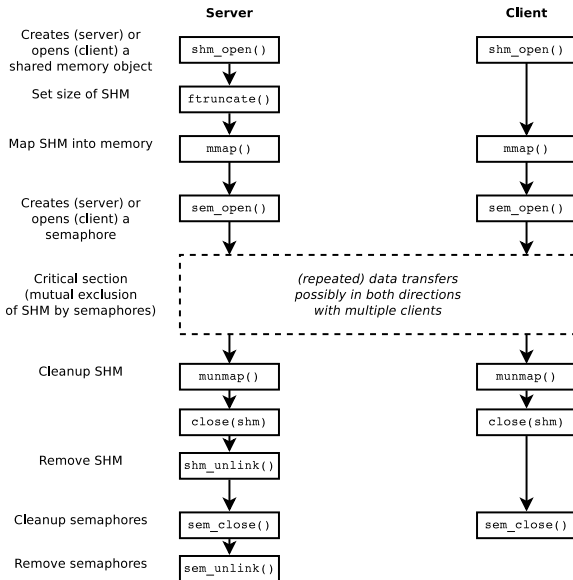
Circular Buffer

Resource Management

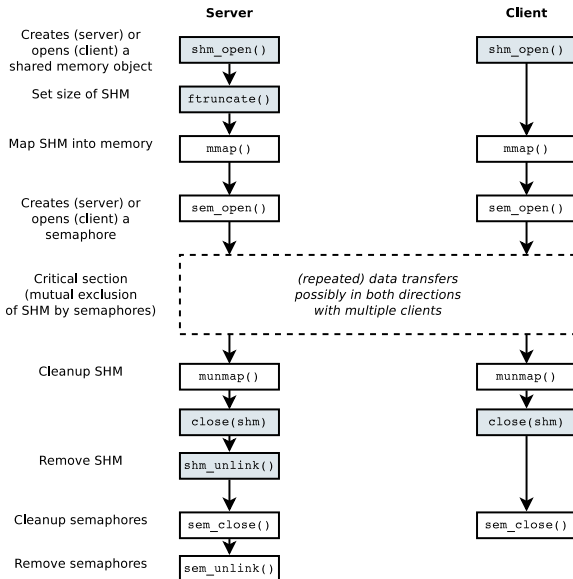
Summary

- ▶ Makes it possible to create shared memory between non-related processes without creating a file
- ▶ **Shared memory objects** identified via names
- ▶ Created on file system for volatile memory: **tmpfs**
- ▶ Behaves as a usual file system (e.g. access rights)
- ▶ Available as long as system is running
- ▶ `mmap` is used to map it into the virtual memory of a process

Client-Server Example



Client-Server Example



Shared Memory API

Create/Open

- ▶ Create and/or open a new/existing object: `shm_open(3)`

```
#include <sys/mman.h>
#include <fcntl.h>      /* For O_* constants */

int shm_open(const char *name, int oflag,
              mode_t mode);
```

name Name like “/somename”

oflag Bit mask: O_RDONLY or O_RDWR and eventually. . .

- ▶ O_CREAT: creates an object unless it is created
- ▶ additionally O_EXCL: error if already created

mode Access rights at creation time, otherwise 0

- ▶ Return value: file descriptor on success,
-1 on error (→ `errno`)
- ▶ Linux: Object at `/dev/shm/somename` created

Shared Memory API

Set Size

- ▶ The creating process normally sets the size (in bytes) based on the file descriptor: `ftruncate(2)`

```
#include <unistd.h>
#include <sys/types.h>

int ftruncate(int fd, off_t length);
```

- ▶ Return value: 0 on success, -1 on error (→ `errno`)
- ▶ Then the file descriptor can be used to create a common mapping (`mmap(2)`) and finally it can be closed (`close(2)`)

Shared Memory API

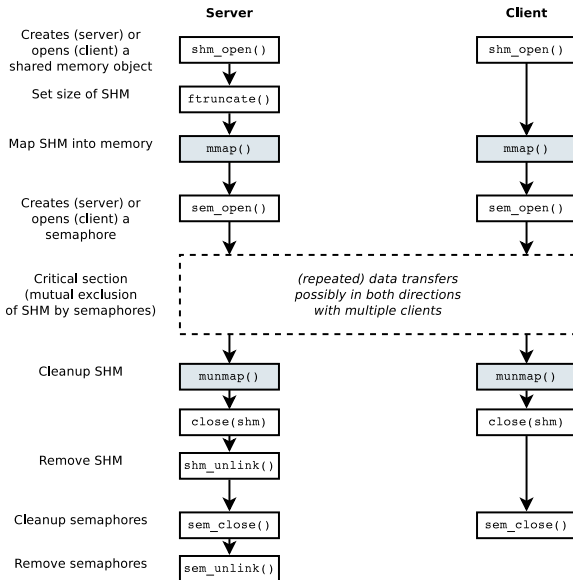
Remove

- ▶ Remove a shared memory object name: `shm_unlink(3)`

```
int shm_unlink(const char *name);
```

- ▶ Name, which was specified at creation
- ▶ Return value: 0 on success, -1 on error (→ `errno`)
- ▶ Further `shm_open()` with the same name raises an error (unless a new object is created by specifying `O_CREAT`)
- ▶ The memory is released when the last process has closed the file descriptor with `close()` and released any mappings with `munmap()`
- ▶ Common commands (`ls`, `rm`) can be used to list and remove `/dev/shm/` (e.g. if program crashes)

Client-Server Example



Memory Mapping

Recall: `mmap(2)`

`mmap(2)`

= maps a file into the virtual memory of a process

- ▶ Multiple processes can access the underlying memory
- ▶ Shared memory is based on sharing a resource (a file)
“shared file mapping”

Memory Mapping

Create

- Create a mapping: `mmap(2)`

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t length, int prot,  
           int flags, int fd, off_t offset);
```

addr Suggestion for starting address, should be NULL

length Size of the mapping in bytes, often the size of a file
(see `fstat(2)`)

prot Bit mask for memory protection: `PROT_NONE` (no
access allowed), `PROT_READ`, `PROT_WRITE`

flags Bit mask, e.g., `MAP_PRIVATE`, `MAP_SHARED`,
`MAP_ANONYMOUS`

fd The file descriptor to be mapped


offset Offset in the file (multiple of page size), 0

- Return value: Starting address of the mapping (aligned to
page limit), `MAP_FAILED` on error (`errno`)

Memory Mapping

Virtual Address Space

virtual memory
of a process



- ▶ Mappings in different processes are created at different addresses
- ▶ Take care by storing pointers!

Memory Mapping

Comments

- ▶ The file descriptor (e.g. of a shared memory) can be closed after the creation of the mapping
- ▶ In Linux, mappings are listed under `/proc/PID/maps`
- ▶ Disadvantages of actual file mappings (not a virtual file) for shared memory: **Persistent** → **costs for disk I/O**
- ▶ For related processes: shared, anonymous mappings (`MAP_SHARED` | `MAP_ANONYMOUS`)
 - ▶ No underlying file, not even a virtual file
 - ▶ Create mapping before `fork()`:
→ child processes can access the mapping at the same address

Memory Mapping

Release

- ▶ Releasing a mapping: `munmap()`

```
#include <sys/mman.h>
```

```
int munmap(void *addr, size_t length);
```

- ▶ Removes whole memory pages from the given space, starting address has to be page-aligned
- ▶ Return value: 0 on success, -1 on error (→ `errno`)

Example

Define Structure of the shared memory

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <unistd.h>

#define SHM_NAME "/myshm"
#define MAX_DATA (50)

struct myshm {
    unsigned int state;
    unsigned int data[MAX_DATA];
};
```


Example

Create and map the shared memory

```
// create and/or open the shared memory object:
int shmfd = shm_open(SHM_NAME, O_RDWR | O_CREAT, 0600);
if (shmfd == -1)
    ... // error

// set the size of the shared memory:
if (ftruncate(shmfd, sizeof(struct myshm)) < 0)
    ... // error

// map shared memory object:
struct myshm *myshm;
myshm = mmap(NULL, sizeof(*myshm), PROT_READ | PROT_WRITE,
             MAP_SHARED, shmfd, 0);

if (myshm == MAP_FAILED)
    ... // error

if (close(shmfd) == -1)
    ... // error
```

Example

Cleanup

Overview

Shared Memory

POSIX
Shared
Memory API

Memory
Mapping

Example

Semaphores

Motivation

Synchronization
Tasks

POSIX
Semaphore
Examples

Circular Buffer

Resource
Management

Summary

```
// unmap shared memory:  
if (munmap(myshm, sizeof(*mysm)) == -1)  
    ... // error  
  
// remove shared memory object:  
if (shm_unlink(SHM_NAME) == -1)  
    ... // error
```

Semaphores

Overview

Shared Memory

POSIX
Shared
Memory API

Memory
Mapping

Example

Semaphores

Motivation

Synchronization
Tasks

POSIX
Semaphore

Examples

Circular Buffer

Resource
Management

Summary

Synchronization

= control access of concurrent processes to a critical section

- ▶ **Conditional synchronization:** In which order is a critical section accessed: A before B? B before A?
- ▶ **Mutual exclusion:** Ensure that only one process is accessing a shared resource ().
Not necessarily fair/alternating.

Example (1)

Thread A:

```
a1: print "yes"
```

Thread B:

```
b1: print "no"
```

- ▶ No deterministic sequence of “yes” and “no”. Depends on, e.g., the scheduler.
- ▶ Multiple calls might cause different outputs. Are other outputs possible?

Example (2)

Thread A:

```
a1: x = 5  
a2: print x
```

Thread B:

```
b1: x = 7
```

- ▶ Path to output “5” and in the end $x = 5$?
- ▶ Path to output “7” and in the end $x = 7$?
- ▶ Path to output “5” and in the end $x = 7$?
- ▶ Path to output “7” and in the end $x = 5$?

Example (3)

Thread A:

```
a1: x = x + 1
```

Thread B:

```
b1: x = x + 1
```

- ▶ Assumption: x is initialized with 1. What are possible values for x after execution?
- ▶ Is $x++$ atomic?

Semaphores

Functions

Semaphore

= “Shared variable” used for synchronization

- ▶ 3 basic operations:
 - ▶ $S = \text{Init}(N)$
create semaphore S with value N
 - ▶ $P(S)$, $\text{Wait}(S)$, $\text{Down}(S)$
decrement S and block when S gets negative
 - ▶ $V(S)$, $\text{Post}(S)$, $\text{Signal}(S)$, $\text{Up}(S)$
increment S and wake up waiting process

Example - Serialization

Overview

Shared Memory

POSIX
Shared
Memory API

Memory
Mapping

Example

Semaphores

Motivation

Synchronization
Tasks

POSIX
Semaphore
Examples

Circular Buffer

Resource
Management

Summary

Thread A:

statement a1

Thread B:

statement b1

How to guarantee that $a1 < b1$ ($a1$ before $b1$)?

Example - Serialization

Initialization:

```
S = Init(0)
```

Thread A:

```
statement a1  
V(S) // post
```

Thread B:

```
P(S) // wait  
statement b1
```

Example - Mutex

Thread A:

$$x = x + 1$$

Thread B:

$$x = x + 1$$

How to guarantee that only one thread is entering the critical section?

Example - Mutex

Initialization:

```
mutex = Init(1)
```

Thread A:

```
P(mutex) // wait  
x = x + 1  
V(mutex) // post
```

Thread B:

```
P(mutex) // wait  
x = x + 1  
V(mutex) // post
```

⇒ Critical section seems to be atomic

Example - Alternating Execution

Thread A:

```
for(;;) {  
    x = x + 1  
}
```

Thread B:

```
for(;;) {  
    x = x + 1  
}
```

How to achieve that A and B are called alternately?

Example - Alternating Execution

Initialization:

```
S1 = Init(1)
S2 = Init(0)
```

Thread A:

```
for(;;) {
    P(S1) // wait
    x = x + 1
    V(S2) // post
}
```

Thread B:

```
for(;;) {
    P(S2) // wait
    x = x + 1
    V(S1) // post
}
```

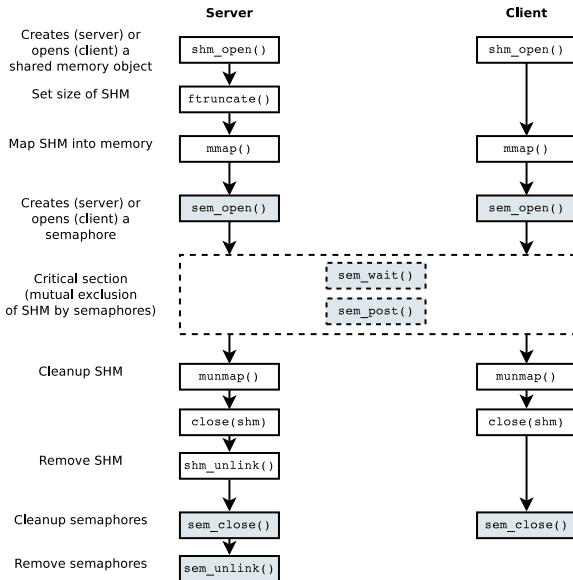
⇒ 2 semaphores are necessary!

How does the synchronization look like for 3 threads that should work alternately? How about N threads?

POSIX Semaphore

- ▶ Synchronization of processes
 - ▶ Non-related processes: [named semaphores](#)
 - ▶ (Related processes or threads within a process: unnamed semaphores)
- ▶ Similar to POSIX shared memory. . .
 - ▶ Identified by name
 - ▶ Created on dedicated file system for volatile memory: [tmpfs](#)
 - ▶ Lifetime limited to system runtime
- ▶ Linked with `-pthread`
- ▶ See also `sem_overview(7)`
- ▶ Linux: object is created at `/dev/shm/sem.somename`

Client-Server Example



Semaphore API

Create/Open

- Create/open a new/existing semaphore: `sem_open(3)`

```
#include <semaphore.h>
#include <fcntl.h>      /* For O_* constants */

/* create a new named semaphore */
sem_t *sem_open(const char *name, int oflag,
                mode_t mode, unsigned int value);

/* open an existing named semaphore */
sem_t *sem_open(const char *name, int oflag);
```

name Name of the form “/somename”

oflag Bit mask: O_CREAT, O_EXCL

mode Access rights (at creation time only)

value Initial value (when creating)

- Return value: **Semaphore address** on success,
SEM_FAILED on error (→ errno)

Semaphore API

Close and Remove

- Close a semaphore: `sem_close(3)`

```
int sem_close(sem_t *sem);
```

- Remove a semaphore: `sem_unlink(3)`

```
int sem_unlink(const char *name);
```

Is released after all processes have closed it.

- Return value: 0 on success, -1 on error (→ `errno`)

Semaphore API

Wait, P()

- ▶ Decrement a semaphore: `sem_wait(3)`

```
int sem_wait(sem_t *sem);
```

- ▶ If the value > 0 , the method returns immediately
- ▶ It blocks the function until the value gets positive otherwise
- ▶ Return value: 0 on success, -1 on error (\rightarrow errno) and the value of the semaphore is not changed

Signal Handling

The function `sem_wait()` can be interrupted by a signal (`errno == EINTR`)!

Semaphore API

Post, V()

- Increment a semaphore: `sem_post(3)`

```
int sem_post(sem_t *sem);
```

- If the value of a semaphore gets positive, a blocked process will continue
- If multiple processes are waiting: the order is not defined (= weak semaphore)
- Return value: 0 on success, -1 on error (\rightarrow errno) and the semaphore value is not changed

Example - Alternating Execution

Process A (code without error handling)

```
#include <stdio.h>
#include <unistd.h>
#include <semaphore.h>
#include <fcntl.h>

#define SEM_1    "/sem_1"
#define SEM_2    "/sem_2"

int main(int argc, char **argv) {
    sem_t *s1 = sem_open(SEM_1, O_CREAT | O_EXCL, 0600, 1);
    sem_t *s2 = sem_open(SEM_2, O_CREAT | O_EXCL, 0600, 0);

    for(int i = 0; i < 3; ++i) {
        sem_wait(s1);
        printf("critical: %s: i = %d\n", argv[0], i);
        sleep(1);
        sem_post(s2);
    }
    sem_close(s1); sem_close(s2);

    return 0;
}
```

Example - Alternating Execution

Process B (code without error handling)

```
#include <stdio.h>
#include <unistd.h>
#include <semaphore.h>
#include <fcntl.h>

#define SEM_1    "/sem_1"
#define SEM_2    "/sem_2"

int main(int argc, char **argv) {
    sem_t *s1 = sem_open(SEM_1, 0);
    sem_t *s2 = sem_open(SEM_2, 0);

    for(int i = 0; i < 3; ++i) {
        sem_wait(s2);
        printf("critical: %s: i = %d\n", argv[0], i);
        sleep(1);
        sem_post(s1);
    }
    sem_close(s1); sem_close(s2);
    sem_unlink(SEM_1); sem_unlink(SEM_2);
    return 0;
}
```

Example - Handling Signals

```
volatile sig_atomic_t quit = 0;

void handle_signal(int signal) { quit = 1; }

int main(void)
{
    sem_t *sem = sem_open(...);

    struct sigaction sa = { .sa_handler = handle_signal; };
    sigaction(SIGINT, &sa, NULL);


    while (!quit) {
        if (sem_wait(sem) == -1) {
            if (errno == EINTR) // interrupted by signal?
                continue;

            error_exit(); // other error
        }

        ...
    }
}
```

Circular Buffer

= simple FIFO implementation with shared memory and semaphores



```
int wr_pos = 0;
void write(int val) {
    sem_wait(free_sem);
    buf[wr_pos] = val;
    sem_post(used_sem);
    wr_pos += 1;
    wr_pos %= sizeof(buf);
}
```

```
int rd_pos = 0;
int read() {
    sem_wait(used_sem);
    int val = buf[rd_pos];
    sem_post(free_sem);
    rd_pos += 1;
    rd_pos %= sizeof(buf);
    return val;
}
```

Circular Buffer (code without error handling)

```
#define BUF_LEN 8
int *buf; // points to shared memory mapped with mmap(2)
sem_t *free_sem, // tracks free space, initialized to BUF_LEN
      *used_sem; // tracks used space, initialized to 0
```

```
int write_pos = 0;
void circ_buf_write(int val) {
    sem_wait(free_sem); // writing requires free space
    buf[write_pos] = val;
    sem_post(used_sem); // space is used by written data
    write_pos = (write_pos + 1) % BUF_LEN;
}
```

```
int read_pos = 0;
int circ_buf_read() {
    sem_wait(used_sem); // reading requires data (used space)
    int val = buf[read_pos];
    sem_post(free_sem); // reading frees up space
    read_pos = (read_pos + 1) % BUF_LEN;
    return val;
}
```

Reading and writing can happen simultaneously!

Resource Management

Who creates a resource?

→ depends on the calling sequence of the processes!

- ▶ Undefined
- ▶ Fixed sequence (e.g., client-server systems)

Who removes resources?

- ▶ No errors during program execution
- ▶ On errors
 - ▶ Unsynchronized cleanup:
Process that has errors removes resources
 - ▶ Synchronized cleanup:
Special communication is necessary (expensive)

Resource Allocation

... at undefined calling sequence

- ▶ Creates it, unless it exists
- ▶ O_CREAT flag without O_EXCL flag (no error if the SHM already exists → SHM will be opened only)
- ▶ e.g., shared memory:

```
shmfd = shm_open(SHM_NAME,  
                 O_CREAT | O_RDWR, PERM);  
  
if (shmfd == -1) ... /* error */
```

Remove Resources

... at fixed calling sequence

Assumption: Synchronization ok (correctly implemented), i.e., processes execute the critical section in the correct sequence

→ Process responsible (the last one accessing the critical section) is able to remove the resources

- ▶ Remove the resources local to the process as usual (e.g., close opened log file)
- ▶ Remove the kernel persistent resources (e.g., shared memory, semaphores) on normal process termination **but also on errors**
- ▶ Help: `atexit(3)`
 - ▶ Register a function that gets called on normal process termination
 - ▶ Multiple functions: in reversed sequence of the registration
 - ▶ Registered functions are not called on `_exit()` (cf. signal handling)

Remove Resources

Example

```
static int shmfd = -1;

void allocate_resources(void) {
    shmfd = shm_open(SHM_NAME, O_CREAT ...);
    ...
}

void free_resources(void) {
    ...
    if (shmfd != -1) {
        if (shm_unlink(SHM_NAME) == -1)
            /* print error message, DON'T CALL EXIT */;
    }
}

void main(void) {
    if (atexit(free_resources) != 0)
        /* error */

    allocate_resources();
    ...
}
```

Summary

Overview

Shared Memory

POSIX
Shared
Memory API
Memory
Mapping
Example

Semaphores

Motivation
Synchronization
Tasks
POSIX
Semaphore
Examples

Circular Buffer

Resource Management

Summary

- ▶ Shared memory is a fast method for IPC
- ▶ Explicit synchronization with semaphores
- ▶ Synchronization tasks
- ▶ Strategies to resource (de-)allocation

Material

Overview

Shared Memory

POSIX
Shared
Memory API

Memory
Mapping

Example

Semaphores

Motivation

Synchronization
Tasks

POSIX
Semaphore

Examples

Circular Buffer

Resource Management

Summary

- ▶ Michael Kerrisk: A Linux and UNIX System Programming Handbook, No Starch Press, 2010.
- ▶ Linux implementation of shared memory/tmpfs:
<http://www.technovelty.org/linux/shared-memory.html>
- ▶ Richard W. Stevens: UNIX Network Programming, Vol. 2: Interprocess Communications